

Live Agent 02 — Opportunity Qualification

Real-time AI bid qualification with multi-reviewer human approval · Built by Deepak Thirukkumaran · June 2026

WHAT THIS IS: The Opportunity Qualification agent is now **live** — every “Run” click sends real tender and bidder data to a large language model and returns a fresh Go / No-Go decision. The mockup demo for the other agents stays intact.

How it works — the user flow

1	Open the agent	Click Live · Real LLM from the sidebar. Lands on the live qualification page.
2	Enter inputs	Two editable fields: Tender Context (the RFP details) and Bidder Profile (your company's strengths). Pre-filled with realistic sample data — edit freely.
3	Run the agent	Click Run Live Agent . The system sends the inputs to a real LLM (Google Gemini 2.5 Flash) and returns a complete analysis in 3–5 seconds.
4	Review the output	Get a Go / No-Go verdict, confidence score, 8-criteria scorecard, strengths, risks, conditions, bid strategy, and likely competitors — all generated live by the AI.
5	Initiate multi-reviewer approval	Click Initiate Review . A chatbot opens and walks you through assigning four reviewers (Technical, Commercial, Legal, Regional) from a team roster.
6	Approvals roll in	Each assigned reviewer can Approve or Request Revision on their lane. Progress bar tracks status. Only when all 4 approve does the agent's output get marked “Approved”.

What the AI generates on each run

✓ Generated live by the LLM • Go / No-Go / Conditional Go verdict • Confidence (0–100%) and overall score • Win probability + estimated contract value • Margin potential range • 8 criterion scores with written rationales • 5 strengths + 5 risks + 3 conditions • Recommended bid strategy paragraph • List of likely competitors with threat levels	■ Pre-defined structure • The 8 scoring categories (Technical, Past Experience, Capacity, Commercial, Risk, Competitive, Strategic, Local Content) • JSON response shape (for consistent rendering) • Default sample inputs (editable by user) <i>These constraints exist so output is consistent and parseable. This is exactly how enterprise AI tools like Copilot and Einstein are built.</i>
--	--

Multi-reviewer human approval — the 4 lanes

Lane	Reviews	Focus Areas
Technical	Engineering capability, capacity, past projects	Technical Capability · Past Experience · Capacity
Commercial	Contract value, margin, commercial terms	Commercial Fit · Strategic Alignment
Legal	Risk clauses, compliance, change-in-law exposure	Risk Exposure · Local Content
Regional	Local market, regulatory, competitive context	Competitive Landscape · Local Content

Each reviewer is assigned via an **Initiate Review chatbot** that walks the user through picking reviewers from a team roster (or typing new names). All 4 must approve before the agent output is marked “Approved”.

How to verify it's a real LLM

Test 1 — Click Run twice with the same inputs. Scores and wording will differ slightly each time. **Test 2** — Edit the bidder profile to something weak (e.g. “a 2-year-old startup with no relevant experience”) and click Run again → the verdict flips to NO-GO with completely different risks and strategy. That's the AI reasoning from your input, not a cached answer.

Currently running on Google Gemini 2.5 Flash (free tier, ~1,500 requests/day). Cost so far: ■0. Switchable to Claude or other LLMs with one config change.