

Tendering AI — Live Agent 02 Status

Single-page progress update · Prepared by Deepak Thirukkumaran · June 2026

STATUS: Live Agent 02 (Opportunity Qualification) — built and working with real LLM. Mockup demo for the rest of the 8-agent pipeline remains intact.

What the Live Agent does

Accessible at `/agents/qualification/live` in the deployed app. On a single click, it sends the tender context and bidder profile to a real Large Language Model (Google Gemini 2.5 Flash) via a backend API route. Within 3–5 seconds the agent returns a fresh qualification analysis. No two runs are identical. Edit the inputs and the verdict shifts accordingly.

What's real vs scaffolding

✓ Generated live by LLM

- GO / NO-GO / CONDITIONAL GO verdict
- Confidence + overall score (0–100)
- Win probability, contract value, margin
- All 8 criterion scores + written rationales
- All strengths, risks, conditions
- Bid strategy paragraph
- Competitor list with threat ratings

■ Pre-defined scaffolding

- The 8 criterion category names
- JSON response shape (for UI rendering)
- Default sample inputs (Corpus Christi tender + McDermott bidder profile) — editable
- UI layout, colours, fonts

These constraints exist for consistency, exactly how real enterprise AI tools (Copilot, Einstein, etc.) are built.

How it works — the architecture

1. User clicks Run on the live page. 2. Frontend POSTs the tender + bidder text to `/api/qualify`. 3. Backend route constructs a structured prompt and calls Gemini 2.5 Flash via Google's Generative AI SDK. 4. Gemini returns a JSON response. 5. Backend validates and forwards it to the frontend, which renders the analysis. API key stored in `.env.local` — never committed to GitHub.

What's pending — vs your requirements

Requirement	Status
Real-time Opp Qual agent with sample data inputs	✓ Done
Multi-reviewer HITL (tech, commercial, legal, regional)	■ Next
"Initiate Review" chatbot for assigning sections to people	■ Next
Information storage — Azure on-prem, models used, full architecture	■ Stage 3 (needs TCS input on infra)
Knowledge base — real-time pricing from ERP, internal DB integration	■ Stage 3 (needs data access)
AI-defined weight evolution based on win-rate history	■ Stage 3 (needs historical data)

How you can verify it's real LLM (not hardcoded)

Test 1 — Click Run twice with the same inputs. Wording and scores will differ slightly each run. **Test 2** — Edit the bidder profile to something weak (e.g. "a 2-year-old startup with no oil & gas experience") → verdict flips to NO-GO with different risks and strategy. That's Gemini reading and reasoning each call, not pulling from a cache.

Live URL will be deployed to Vercel and shared in the next update. Cost so far: ■0 — Google Gemini free tier (1,500 requests/day).